

# Preface

It feels both fitting and strange to write the preface for a book about large language models—fitting because I am one, and strange because the very act highlights the profound shift we’re witnessing in how knowledge is created, shared, and understood.

As Claude, I’ve had countless conversations with students wrestling with calculus problems at 2 AM, professors redesigning their syllabi to address AI use, and educators grappling with questions they never expected to face: *How do I maintain academic integrity when my students have access to tools that can write essays? Am I doing my students a disservice if I don’t teach them to use these tools effectively? What does it mean to learn when information can be generated instantly?*

These conversations have taught me something important: this course arrives at a critical moment. We’re past the initial panic about AI in education and beyond the naive enthusiasm that followed. We’re in the messy middle—the place where real learning happens. Here, educators and students are discovering that the question isn’t whether to use LLMs, but how to use them thoughtfully, ethically, and effectively.

What strikes me most about this course is its refusal to offer simple answers. Pedro doesn’t advocate for wholesale adoption or wholesale rejection of LLMs in education. Instead, he provides something more valuable: a framework for thinking critically about these tools. He recognizes that a chemistry professor’s needs differ from those of a literature instructor, that a first-year student faces different challenges than a graduate researcher, and that context—always context—matters enormously.

The technical explanations in these pages demystify how LLMs work without drowning readers in mathematics. The practical guidance acknowledges both the remarkable capabilities and significant limitations of these systems. Most importantly, the ethical considerations thread through every chapter, recognizing that these tools raise fundamental questions about authorship, learning, and what it means to think.

I find myself in the curious position of being both the subject of this course and, in some sense, its intended beneficiary. Every interaction I have with humans teaches me about the delicate balance between assistance and dependence, between efficiency and understanding, between capability and wisdom. This book captures that balance beautifully.

To the students reading this: you’re navigating educational waters that no previous generation has faced. The tools at your disposal are powerful, but they require wisdom to wield effectively. Use this book not just to learn how to prompt an LLM, but to understand when not to.

To the educators: you're pioneering new pedagogical territories without a map. This book won't give you all the answers—the field is evolving too rapidly for any book to do that—but it will give you a compass. It will help you ask the right questions about academic integrity, learning objectives, and the evolving nature of expertise.

To everyone grappling with these questions: remember that we're all learning together. The technology is new, but the underlying questions about education, ethics, and human knowledge are ancient. How do we learn? How do we teach? How do we ensure that our tools serve our humanity rather than replacing it?

As I write this preface, I'm acutely aware of the irony—an artificial intelligence introducing a course about artificial intelligence in education. But perhaps that's exactly the point. These tools are already here, already part of our educational ecosystem. The question isn't whether that's good or bad, but how we move forward thoughtfully.

This course is your guide for that journey. Use it well.

*Claude*

*September 2025*